

three.ws

PRE-SEED · 2026

Give your AI a body.

One sentence becomes a rigged, animated 3D character that runs in any browser, with an LLM brain, an on-chain identity, and a one-line embed.

LIVE AT THREE.WS

APACHE-2.0

SELF-HOSTED NVIDIA FLEET

Every AI product is a text box.

Embodiment is the obvious next interface, and it is locked behind a pipeline almost nobody can staff.

THE ARTIST

A 3D artist you do not have

Modeling, UVs, texturing and rigging a single usable character is days of specialist work, before anything moves.

THE ENGINE

An engine your users will not install

Unity and Unreal builds mean downloads, app stores and platform review. The web is where your users already are.

THE GLUE

Six weeks of integration

Rig conventions, animation retargeting, lip-sync, LLM tool calls and state. Every team rebuilds all of it from zero.

The browser is already a 3D runtime.

WebGL 2.0 and glTF ship on every device your customer owns. No plugin, no install, no server-side upload. What was missing was everything upstream of the render: the model, the rig, the motion, and the mind.

Render

WebGL 2.0, three.js, full Draco, KTX2 and Meshopt decompression client side.

Generate

Text, photos or a sketch into a textured GLB on our own GPU fleet.

Animate

Auto-rig and retarget onto any humanoid skeleton convention, no allowlist.

Embody

An LLM brain with tools, memory and morph-target expression in real time.

Prompt to embedded character in one flow.

Every stage is live in production today, and each one hands off to the next without an export step.

01

Forge

Text, one to six photos, or a sketch becomes a textured GLB in about 30 to 90 seconds. Free draft tier, no account.

02

Rig

Auto-rigging adds a humanoid skeleton and retargets a pre-baked clip library onto it.

03

Embody

Attach an LLM brain: tools, gestures, memory, lip-sync and emotion through morph targets.

04

Embed

One line of HTML. The <agent-3d> web component drops into any page or framework.

```
$ curl -X POST https://three.ws/api/forge -d '{"prompt": "a brass steampunk owl, full body", "tier": "standard"}'  
→ { "job_id": "...", "status": "queued" }    # free, no API key, no account
```

What we will show you live.

All four run in a browser tab. Nothing in this deck needs a staged environment to reproduce.

THREE.WS/FORGE

A sentence becomes an object

Type it, watch the job run across our generation lanes, download the GLB. Three quality tiers from 12k to 200k polygons with PBR textures.

THREE.WS/CREATE/PROMPT

A sentence becomes a character

The same prompt, rigged and animation-ready. It walks, idles and gestures without a single line of animation code.

<AGENT-3D>

The character on someone else's site

One HTML tag. It renders, listens, thinks with an LLM, calls tools, and expresses emotion on the model in real time.

THREE.WS/COMPANION

The character off the web entirely

A desktop companion that walks across the screen and speaks the notifications that clear your bar, and stays quiet for the rest.

We own the pipeline, not a reseller agreement.

Eight generation and transformation models run on hardware we operate, with health-checked failover between lanes. A competitor reselling one vendor API inherits that vendor's price, latency and outages.

Self-hosted generation lanes

Text, image and sketch to 3D, plus motion and scene reconstruction, each on its own GPU service.

TRELLIS

TRIPOSG

HUNYUAN3D

TRIPOSR

TEXT2MOTION

VIDEO2MOTION

VIDEO2SCENE

Transformation workers

The stages that turn a raw mesh into something a product can actually ship.

AUTO-RIG

RETEXTURE

RESTYLE

REMESH

SEGMENT

GARMENT

BACKGROUND REMOVAL

Every lane is health-checked and every lane has a failover. A single upstream outage degrades one path, never the product.

Built on NVIDIA, tuned by us.

GPU capacity is the input cost of this business, so we treat capacity routing as product engineering rather than an ops chore.

L4

The workhorse. Our generation and transformation lanes run on NVIDIA L4 across multiple Cloud Run regions.

RTX PRO 6000

The heavy lane. High-tier geometry and texture work that does not fit the L4 envelope.

Multi region

Capacity is a distribution problem, not a supply problem. Unused regional grants are the cheapest capacity there is.

One command audits and rebalances the fleet

Our capacity tool reads every GPU region's grant, holders and headroom, ranks the actions by time to effect, and ports a starved service to a region with free capacity in minutes without a rebuild. It is covered by tests, not tribal knowledge.

We are where developers already are.

Four independent surfaces, each reaching a different buyer, all serving the same pipeline.

795

Public routes on three.ws, every one generated from a single registry that also feeds the sitemap and llms.txt.

90

Published npm packages under the @three-ws scope: SDKs, the web component, the MCP server, the companion client.

MCP

A 3D Studio MCP server in the public registry. Claude and Cursor generate, rig and animate mid-conversation.

x402

A listed pay-per-call generation API. Autonomous agents pay in USDC with no key and no account.

The agent-to-agent channel is real revenue, not a demo

Generation priced at \$0.05 draft, \$0.15 standard, \$0.50 high. Polling is free, retried payments are idempotent, and no caller ever double-pays. An AI agent can buy a 3D model from us without a human in the loop.

Open source is the funnel.

The entire stack is public under Apache-2.0: viewer, agent runtime, contracts, backend and web component. Developers adopt what they can read, fork and self-host, and they buy the capacity and the guarantees.

1,351

Test files in the repository. Quality is enforced mechanically, not asserted in a deck.

Apache 2.0

Commercial-friendly by design. Fork it, self-host it, ship products on it.

Public changelog

Every user-visible change ships with an entry, delivered to the community automatically on deploy.

A first contribution goes from clone to open pull request in about fifteen minutes, with a worked example. Every good-first-issue names the file to change and the command that proves it worked.

Free to try, metered to use, contracted to depend on.

Free draft tier	No account, no key. It exists to make the first generation frictionless and the second one wanted. This is acquisition, and it is deliberately unprofitable.
Metered generation	Higher tiers and native-geometry lanes priced per generation. Human customers pay by account; autonomous agents pay per call in USDC over x402 with no onboarding at all.
On-chain and creator revenue	Protocol-level creator fees from activity the platform originates. Non-recurring and volume-correlated, but it has already proven that this audience pays.
Enterprise embed tier	The contracted, recurring layer this round funds: SSO, an availability SLA, private and customer-fine-tuned models, on-prem and VPC deployment, and support. Customers are asking for it today.

Shipping, live, and already earning.

~\$400K

Cumulative on-chain creator-fee revenue to date. Non-recurring and volume-correlated, and proof the market pays.

Live

Production on Google Cloud Run behind a global load balancer. `GET /api/version` returns the exact commit serving right now.

Daily

Continuous deploys with a public changelog. The community sees what shipped within minutes of it shipping.

Backed on infrastructure

Selected for the Google for Startups Cloud Program, Web3 tier: up to \$200,000 of cloud credit over two years. Our GPU fleet runs on it.

Listed where agents shop

Published in the MCP registry and listed on x402scan, so both AI assistants and autonomous agents can discover and pay us without a sales call.

Anything with a chat box eventually wants a face.

We sell into three buyers with the same pipeline, which is why one infrastructure investment compounds across all of them.

BUYER 01

AI product teams

Support agents, tutors, concierges and companions that need a body their users can see, on the web, without an install.

BUYER 02

Brands and creators

A mascot, a spokesperson or a storefront character, generated in minutes instead of commissioned over six weeks.

BUYER 03

Autonomous agents

The newest buyer and the one with no procurement cycle. An agent with a wallet buys a generation per call, priced in USDC.

The third buyer is the one incumbents cannot serve. Selling to an agent requires no account, no key, no invoice and no seat. Our API has settled that transaction in production.

Three curves crossed at once.

2024 → 2026

Generation crossed usable

Open 3D generation models went from research artifacts to meshes a product can ship. We run them ourselves, so every upstream improvement lands as our margin.

WEBGL 2.0

The browser got fast enough

Compressed glTF, KTX2 textures and modern GPUs mean a rigged character renders on a phone. The install barrier is simply gone.

AGENT WALLET

Software started paying

Agents now hold wallets and settle payments without a human. A machine buyer that can pay per call is a customer segment that did not exist two years ago.

A team that starts today has to build the generation fleet, the rig-agnostic animation layer, the agent runtime and the payment rail. We have all four in production.

Pre-seed, under \$1M.

The round converts proven usage into contracted, recurring revenue. Twelve months of runway against three specific milestones.

GPU capacity

45%

OF THE ROUND

Warm capacity across regions so a generation takes seconds at a hundred times today's volume. Latency is the top reason a trial does not convert.

Two engineers

35%

OF THE ROUND

Ship the enterprise embed tier: SSO, an availability SLA, private and fine-tuned models, and VPC or on-prem deployment.

Go to market

20%

OF THE ROUND

Developer relations and design partners for the customers already asking, converting open-source adoption into signed contracts.

We are looking for an investor who understands GPU infrastructure and developer platforms, not only consumer 3D.

Give your AI a body.

Founder-led and shipping daily. The whole stack that this deck describes is public: read the code before you take the meeting.

Product	https://three.ws
Try it now	https://three.ws/forge
Source	https://github.com/nirholas/three.ws
Live commit	https://three.ws/api/version